

Uncertainty Quantification for Satellite Image Segmentation

Mohamed Hasan

May 15, 2024

1 Introduction

This report details the development and evaluation of a custom Artificial Neural Network (ANN) designed for image segmentation tasks. I focus on a UNet model architecture, built from scratch using TensorFlow 2, to perform segmentation on satellite images. To enhance the robustness and reliability of the model, I implement Monte Carlo (MC) Dropout as a method for uncertainty quantification. This approach allows me to assess the confidence of the model's predictions, providing valuable insights into the reliability of segmentation results.

The experiments involve an evaluation of different activation functions, loss functions, and other hyperparameters, with a focus on their impact on uncertainty quantification. The satellite image dataset, comprising of a mix of real and synthetic images and corresponding masks, serves as a challenging real-world scenario due to factors like variable lighting conditions, diverse backgrounds, and the presence of objects (i.e. asteroids).

2 Dataset

The dataset consists of satellite images and corresponding masks of size 1280×720 , with each mask delineating the satellite from the background. The images are of varying resolutions and quality, reflecting the diversity of satellite imagery encountered in real-world scenarios. The dataset is divided into training and validation sets, with the training set containing 3117 images and the validation set comprising 600 images. The masks are categorized into fine and coarse masks, with 403 fine masks and 2114 coarse masks in the training set. The dataset presents a challenging segmentation task due to the presence of complex textures, varying lighting conditions, and the need to accurately delineate the satellite from the background.




Figure 1: Sample image and mask from the dataset. Image size is 1280×720 .

2.1 Data Preprocessing

The dataset undergoes preprocessing to ensure uniformity and compatibility with the model. The images are resized to a fixed resolution of 256×256 pixels and The masks are converted to binary format, with pixel values of 0 and 1 representing the background and satellite, respectively, to reduce computational complexity. The preprocessed dataset has the same split between training and validation sets as the original dataset.

The file `data_loader_unit.py` in the `utils` folder contains the function `load_and_process_files()` to load the dataset and preprocess the images and masks. After processing, the data will be saved as numpy arrays for easy access during training. The saved files are named as follows:

- **prepped_data/trainimages.npy**: Contains the preprocessed training images.
- **prepped_data/trainmasks.npy**: Contains the preprocessed training masks.
- **prepped_data/valimages.npy**: Contains the preprocessed validation images.
- **prepped_data/valmasks.npy**: Contains the preprocessed validation masks.

The function to check if the prepped files exist is `check_prepped_data()` in the `utils` folder. It is used in the scripts `train.py`, `predict.py`, and `main.py`. If the prepped files do not exist, the function is called to create them. Now, let's look at the processed dataset.




Figure 2: Sample processed image and mask from the dataset. Image size is 256×256 now and the mask is binary.

3 Model Architecture

The model architecture is based on the UNet architecture, a popular choice for image segmentation tasks due to its ability to capture both local and global features effectively. The model consists of an encoder and a decoder, with skip connections between corresponding layers to preserve spatial information. The encoder downsamples the input image to extract features, while the decoder upsamples the features to generate the final segmentation map. The model is implemented using TensorFlow 2 and Keras, further implementation details can be found in the model definition.

The model architecture is defined in the file `unet.py` in the `model` folder. The implementation is nearly identical to the standard UNet architecture. We add dropout layers after each convolutional layer in the encoder and decoder, so that we may experiment with Monte Carlo Dropout for uncertainty quantification.

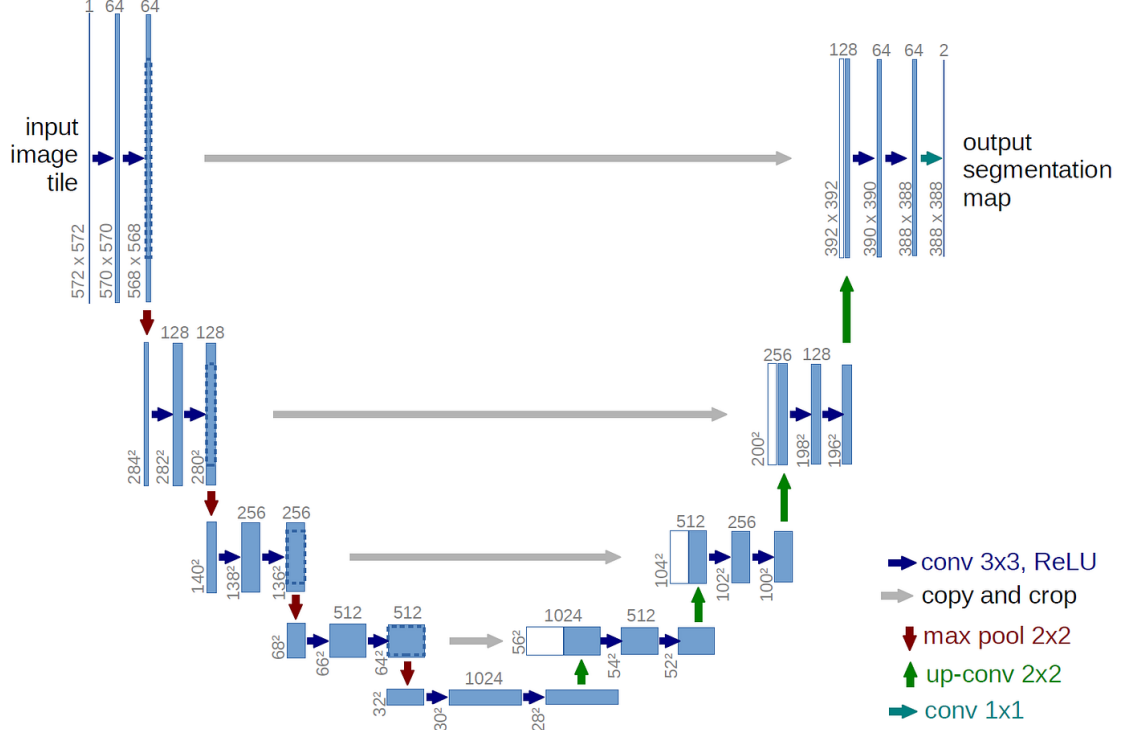


Figure 3: Original UNet model architecture. We use a similar architecture with dropout layers for uncertainty quantification.

3.1 Loss Function, Metrics, and Hyperparameters

The loss function used for training the model is a combination of Binary Cross Entropy (BCE) and Dice Loss. The BCE loss is effective in handling class imbalance, while the Dice Loss directly optimizes for the Dice Coefficient, a common metric in image segmentation tasks. The Dice Coefficient is defined as

$$\text{Dice Coefficient} = \frac{2 \times \text{Intersection}}{\text{Union} + \text{Intersection}} = \frac{2 \sum_{i=1}^N y_i p_i}{\sum_{i=1}^N y_i + \sum_{i=1}^N p_i},$$

$$\text{Dice Loss} = 1 - \text{Dice Coefficient},$$

Where y_i is the ground truth label, p_i is the predicted probability, and N is the total number of pixels. The intersection is the number of overlapping pixels between the predicted and ground truth masks, and the union is the total number of pixels in both masks. The BCE loss is given by

$$\text{BCE} = -\frac{1}{N} \sum_{i=1}^N [y_i \log(p_i) + (1 - y_i) \log(1 - p_i)],$$

A combined loss function is used to leverage the strengths of both BCE and Dice Loss.

$$\text{Combined Loss} = \text{BCE} + \text{Dice Loss}.$$

The model is trained using the Adam optimizer with a learning rate of 0.001. The metrics used for evaluation are the Dice Coefficient and Accuracy.

$$\text{Accuracy} = \frac{\text{Number of Correct Predictions}}{\text{Total Number of Predictions}}$$

The hyperparameters and model configurations are defined in the file `config.py` in the `model` folder. The file contains the following hyperparameters:

- **BATCH_SIZE** = 1: The batch size used for training the model.
- **EPOCHS** = 5: The number of epochs for training.
- **DROPOUT_RATE** = 0.35: The dropout rate used in the model.

3.2 Evaluation

The final model achieved:

- Loss: 0.399
- Accuracy: 0.940
- Dice Coefficient: 0.803

The model shows good performance on the validation set, with high accuracy and Dice Coefficient, despite the challenging nature of the dataset and the low number of epochs. Let's visualize the model's predictions on a sample image from the validation set.

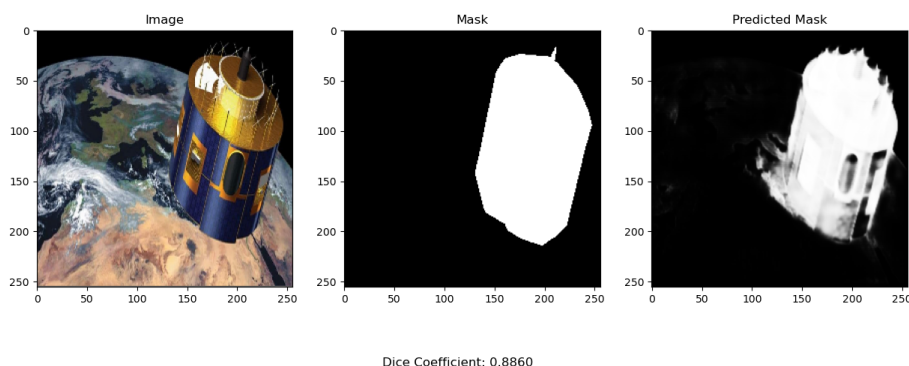


Figure 4: Sample prediction from the validation set.

The model successfully segments the satellite from the background, capturing the main features of the satellite with a dice coefficient of 0.8860. Notice however that the model captures the details of the satellite which we did not include in the masks. This is due to the model's ability to learn from the training data and generalize.

4 Uncertainty Quantification

4.1 Monte Carlo Dropout

For uncertainty quantification, I implemented the Monte Carlo (MC) Dropout technique. Dropout is a regularization method used during the training of neural networks to prevent overfitting. It works by randomly "dropping out" a subset of neurons during each forward pass, effectively creating an ensemble of different network architectures. Typically, dropout is disabled during inference to use the full capacity of the

trained model. However, in MC Dropout, dropout is kept active during inference to introduce stochasticity into the model's predictions.

By incorporating dropout layers with a dropout rate of 0.35 during inference and performing multiple forward passes (in this case, 10 passes), I generated a distribution of predictions for each input image. This approach allowed me to capture the variability in the model's predictions due to the stochastic dropout, which can be interpreted as a measure of uncertainty.

The mean prediction for each pixel is calculated as follows:

$$\mu(x) = \frac{1}{T} \sum_{t=1}^T \hat{y}_t(x),$$

where $\hat{y}_t(x)$ is the prediction at the t -th forward pass and T is the number of passes. The mean prediction provides an aggregate view of the model's output across multiple passes.

The uncertainty is quantified as the standard deviation of these predictions:

$$\sigma(x) = \sqrt{\frac{1}{T} \sum_{t=1}^T (\hat{y}_t(x) - \mu(x))^2}.$$

This standard deviation reflects the model's confidence in its predictions. A higher standard deviation indicates greater uncertainty, whereas a lower standard deviation suggests higher confidence.

MC Dropout is particularly useful for identifying regions where the model is less certain about its predictions. In the context of image segmentation, uncertainty is often higher around the edges of objects and in regions with complex textures. This can be seen in Figure 5, where we can see the mean prediction and the standard deviation uncertainty generated by the MC Dropout technique for one of the validation images.

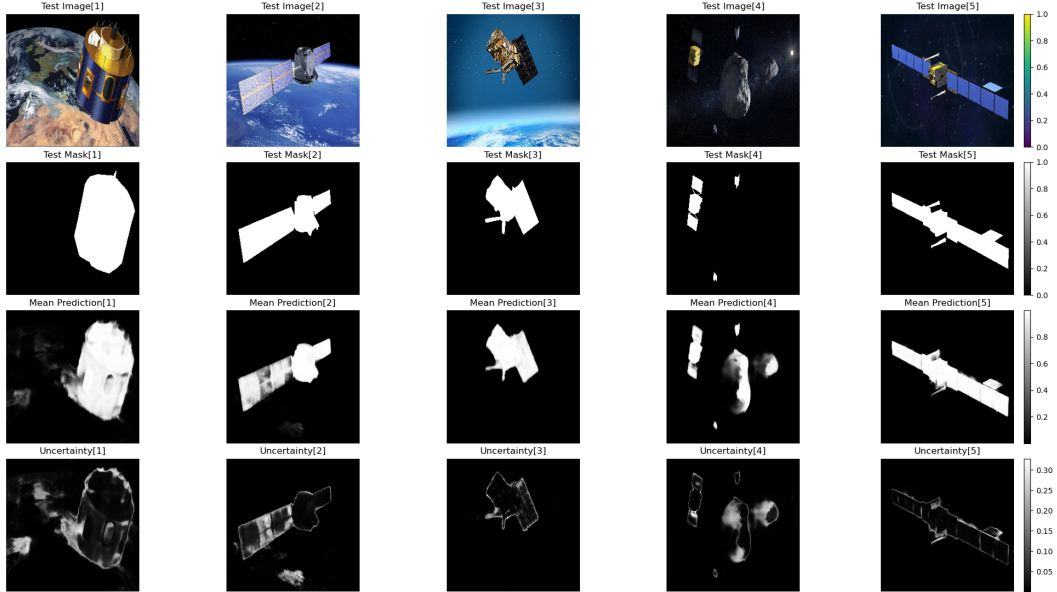


Figure 5: We can see the test image, its mask, the mean prediction, and the uncertainty.

- **test image[1]:** The satellite's body is well captured, but the model shows high uncertainty around the edges and finer details. We can also see that the uncertainty makes sense, especially near the bottom lower edge of the satellite where there is a body of water. The water and the satellite are both blue in this case, so the model is less certain about the prediction in that region.

- **test image[2]**: The satellite and its solar panels are identified, but the prediction is somewhat spread, with significant uncertainty around the solar panels. This is most likely because the solar panels are blue and the background is the blue oceans which makes it harder for the model to distinguish between the two.
- **test image[3]**: The satellite is identified quite well, with uncertainty only around the edges.
- **test image[4]**: The three satellites are identified, but the model also identifies two asteroids as satellites. This is most likely due to the dataset not having many examples of asteroids and the asteroids having similar features (most likely textures) to the satellites.
- **test image[5]**: The satellite is identified quite well, with uncertainty only around the edges.

4.2 Activation Functions

Activation functions play a crucial role in the learning and representation capabilities of neural networks. They introduce non-linearity, enabling the network to learn complex patterns. Each activation function has its unique characteristics and impacts the model's convergence and generalization in different ways. By evaluating multiple activation functions, we aim to identify which ones contribute most effectively to the accuracy and robustness of our segmentation model.

| Activation Function | Loss | Accuracy | Dice Coefficient |
|---------------------|-------|----------|------------------|
| ReLU | 0.399 | 0.94 | 0.803 |
| ELU | 0.499 | 0.914 | 0.773 |
| Swish | 0.43 | 0.934 | 0.762 |
| GELU | 0.436 | 0.932 | 0.775 |
| Leaky ReLU | 0.418 | 0.933 | 0.772 |

Table 1: Performance Metrics for Different Activation Functions

From the results in Table 1, we observe that the ReLU activation function performs the best in terms of accuracy and Dice. Good thing we used ReLU as the activation function in our model for the MC Dropout uncertainty quantification, as it is the best performing activation function in this case.